

Nzengou E. D.

v21eugnz@du.se

```
pd.to_datetime('today')
```

```
Timestamp('2021-04-13 14:17:13.825035')
```

▼ AMI23B – Business Intelligence Lab1

▼ Using Pandas to Explore a Dataset

Task1: Setting up the environment

Using colab, with all the dependencies already installed.

▼ Task2: Using the Pandas Python Library

▼ 2.1) Create a download script download_nba_all_elo.py to download the data

```
import requests

#get the data from a download link
download_url = "https://raw.githubusercontent.com/fivethirtyeight/data/master/nba-elo/nbaallelo.csv"

target_csv_path = "nba_all_elo.csv"

response = requests.get(download_url)
response.raise_for_status() #check that the request was successful

#save the file nba_all_elo.csv in your current working directory.
with open(target_csv_path, "wb") as f:
    f.write(response.content)
    print("download ready")

download ready
```

- ▼ 2.2) Now create a new script lab2_NBA in which you will use the Pandas Python library to take a look at your data

```
#importing Pandas in Python with the pd alias
import pandas as pd

#read in the dataset and store it as a DataFrame object in the variable nba

nba = pd.read_csv("nba_all_elo.csv")

#check nba's type, it should be a DataFrame
type(nba)

pandas.core.frame.DataFrame
```

- ▼ 2.3) Let's see how much data is actually in nba (report these findings)

```
#len() determines the number of rows (observations) in a dataset
len(nba)
```

```
126314
```

```
#.shape determines dimensionality
#the result is a tuple containing number of rows and columns
nba.shape
```

```
(126314, 23)
```

```
#take a look at the first five rows to see the actual data
nba.head()
```

	gameorder	game_id	lg_id	_iscopy	year_id	date_game	seasongame	is_playoff
0	1	194611010TRH	NBA	0	1947	11/1/1946	1	1
1	1	194611010TRH	NBA	1	1947	11/1/1946	1	1
2	2	194611020CHS	NBA	0	1947	11/2/1946	1	1
3	2	194611020CHS	NBA	1	1947	11/2/1946	2	1
4	3	194611020DTF	NBA	0	1947	11/2/1946	1	1

```
#configure Pandas to display all 23 columns
pd.set_option("display.max.columns", None)
```

```
#show only two decimal places
pd.set_option("display.precision", 2)
```

```
#display last five rows
nba.tail()
```

	gameorder	game_id	lg_id	_iscopy	year_id	date_game	seasongame	is_p
126309	63155	201506110CLE	NBA	0	2015	6/11/2015	100	
126310	63156	201506140GSW	NBA	0	2015	6/14/2015	102	
126311	63156	201506140GSW	NBA	1	2015	6/14/2015	101	
126312	63157	201506170CLE	NBA	0	2015	6/16/2015	102	
126313	63157	201506170CLE	NBA	1	2015	6/16/2015	103	

Question.1:

Display the first 3 rows of your dataset.
Remember that the default of `nba.head()` shows the first 5 rows.

```
nba.head(3)
```

	gameorder	game_id	lg_id	_iscopy	year_id	date_game	seasongame	is_playoff
0	1	194611010TRH	NBA	0	1947	11/1/1946	1	
1	1	194611010TRH	NBA	1	1947	11/1/1946	1	
2	2	194611020CHS	NBA	0	1947	11/2/1946	1	

▼ Task3: Get to Know Your Data.

▼ 3.1) Discover the different data types your dataset contains.

```
#display all columns and their data types with .info()
nba.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 126314 entries, 0 to 126313
Data columns (total 23 columns):
```

#	Column	Non-Null	Count	Dtype
0	gameorder	126314	non-null	int64
1	game_id	126314	non-null	object
2	lg_id	126314	non-null	object
3	_iscopy	126314	non-null	int64
4	year_id	126314	non-null	int64
5	date_game	126314	non-null	object
6	seasongame	126314	non-null	int64
7	is_playoffs	126314	non-null	int64
8	team_id	126314	non-null	object
9	fran_id	126314	non-null	object
10	pts	126314	non-null	int64
11	elo_i	126314	non-null	float64
12	elo_n	126314	non-null	float64
13	win_equiv	126314	non-null	float64
14	opp_id	126314	non-null	object
15	opp_fran	126314	non-null	object
16	opp_pts	126314	non-null	int64
17	opp_elo_i	126314	non-null	float64
18	opp_elo_n	126314	non-null	float64
19	game_location	126314	non-null	object
20	game_result	126314	non-null	object
21	forecast	126314	non-null	float64
22	notes	5424	non-null	object

dtypes: float64(6), int64(7), object(10)
memory usage: 22.2+ MB

▼ 3.2) Showing basic statistics

Get an idea of the values each column contains.

```
#get an idea of the values each column contains
# by showing basic descrip
nba.describe()
```

```
# provide other data types by using the include parameter
import numpy as np
```

```
nba.describe(include=np.object)
```

	game_id	lg_id	date_game	team_id	fran_id	opp_id	opp_fran	game_locat
count	126314	126314	126314	126314	126314	126314	126314	126314
unique	63157	2	12426	104	53	104	53	126314
top	198204160ATL	NBA	4/13/2011	BOS	Lakers	BOS	Lakers	126314

Question.2:

Take a look at the team_id and fran_id (franchise) columns, what observations can you make at this point?

At first sight, it appears that team and franchise do not match: the dataset contains almost twice as many unique team IDs (104) as unique franchise IDs (53). This is surprising, we would have expected to have same number for both sides. Additionally, the most common ID is BOS for team, Lakers for franchise. This requires deeper exploration.

▼ 3.3) Exploring the dataset

Exploratory data analysis helps answer questions about your dataset, for example, exploring how often specific values occur in a column. Let's take a look at the two columns team_id and fran_id.

```
nba["team_id"].value_counts()
```

```
BOS    5997
NYK    5769
LAL    5078
DET    4985
PHI    4533
...
PIT     60
INJ     60
DTF     60
TRH     60
SDS     11
Name: team_id, Length: 104, dtype: int64
```

```
nba["fran_id"].value_counts()
```

Lakers	6024
Celtics	5997
Knicks	5769
Warriors	5657
Pistons	5650
Sixers	5644
Hawks	5572
Kings	5475
Wizards	4582
Spurs	4309
Bulls	4307
Pacers	4227
Thunder	4178
Rockets	4154
Nuggets	4120
Nets	4106
Suns	4080
Bucks	4034
Trailblazers	3870
Cavaliers	3810
Clippers	3733
Jazz	3555
Mavericks	3013
Heat	2371
Pelicans	2254
Magic	2207
Timberwolves	2131
Grizzlies	1657
Raptors	1634
Hornets	894
Colonels	846
Squires	799
Spirits	777
Stars	756
Sounds	697
Baltimore	467
Floridians	440
Condors	430
Capitols	291
Olympians	282
Sails	274
Stags	260
Bombers	249
Steamrollers	168
Packers	72
Redskins	65
Rebels	63
Denver	62
Waterloo	62
Huskies	60
Falcons	60
Ironmen	60
Jets	60

```
Name: fran_id, dtype: int64
```

It seems that a team named "Lakers" played 6024 games, but only 5078 of those were played by the Los Angeles Lakers. To find out who the other "Lakers" team is execute the following line of code

```
nba.loc[nba["team_id"] == "Lakers", "team_id"].value_counts()
```

```
LAL    5078
MNL     946
Name: team_id, dtype: int64
```

pandas.DataFrame.loc is used to access a group of rows and columns by label(s) or a boolean array. The output shows that the Minneapolis Lakers ("MNL") played the other 946 games. Let's find out when they played those games:

```
#Find out when they played those games
nba.loc[nba["team_id"] == "MNL", "date_game"].min()
#aggregate the two functions
```

```
'1/1/1949'
```

```
nba.loc[nba["team_id"] == "MNL", "date_game"].max()
```

```
'4/9/1959'
```

```
nba.loc[nba["team_id"] == "MNL", "date_game"].agg(("min", "max"))
```

```
min    1/1/1949
max    4/9/1959
Name: date_game, dtype: object
```

It looks like the Minneapolis Lakers played between the years of 1949 and 1959. That explains why you might not recognize this team!

Question.3 (report your answer):

Find out how many wins and losses the Minneapolis Lakers had, also find how many points they scored

```
nba.loc[nba["team_id"] == "MNL", "game_result"].value_counts()
```

```
W    524
L    422
Name: game_result, dtype: int64
```

Minneapolis Lakers won 524 times, lost 422 times (between 1949 and 1959).

```
mn1_scores = nba.loc[nba["team_id"] == "MNL", "pts"].value_counts()
```

```
mn1_scores.tail(3)
```

```
125    1
127    1
18     1
Name: pts, dtype: int64
```

```
mn1_scores.shape
```

```
(81,)
```

```
# total pts over all MNL match
nba.loc[nba["team_id"] == "MNL", "pts"].sum()

88229
```

Minneapolis Lakers have scored 626484 points during all matches contained in this dataset.

```
# save in csv format
mn1_scores.to_csv("mn1_scores.csv")
```

Question.4:

Now you understand why the Boston Celtics team "BOS" played the most games in the dataset, find out how many points they scored.

```
# total pts over all BOS match
nba.loc[nba["team_id"] == "BOS", "pts"].sum()

626484
```

Boston Celtics have scored 626484 points during all matches contained in this dataset.

Question.5:

After having explored your dataset, explain your observations from Question.2 in a structured way

▼ Task4: Data access methods (loc and iloc):

Check Pandas official docs for these two functions. With data access methods like .loc and .iloc, you can select just the right subset of your DataFrame to help you answer questions about your dataset. .loc uses the label and .iloc the positional index

Question.6 (report your answer):

6.1) Use a data access method to display the 4th row from the bottom of the nba dataset.

```
nba.iloc[-4]
```

```
gameorder      63156
game_id        201506140GSW
lg_id           NBA
_iscopy         0
year_id         2015
date_game       6/14/2015
seasongame      102
is_playoffs     1
team_id         GSW
fran_id         Warriors
pts            104
elo_i           1.8e+03
elo_n           1.8e+03
win_equiv       68
opp_id          CLE
opp_fran        Cavaliers
opp_pts         91
opp_elo_i       1.7e+03
opp_elo_n       1.7e+03
game_location   H
game_result     W
forecast        0.77
notes           NaN
Name: 126310, dtype: object
```

6.2) Use a data access method to display the 2nd row from the top of the nba dataset.

```
nba.iloc[2]
```

```

gameorder      2
game_id        194611020CHS
lg_id          NBA
_iscopy        0
year_id        1947
date_game      11/2/1946
seasongame     1
is_playoffs    0
team_id        CHS
fran_id       Stags
pts            63
elo_i          1.3e+03
elo_n          1.3e+03
win_equiv      42
opp_id         NYK
opp_fran       Knicks
opp_pts        47
opp_elo_i      1.3e+03
opp_elo_n      1.3e+03
game_location  H
game_result    W
forecast       0.63
notes         NaN
Name: 2, dtype: object

```

6.3) Access all games between the labels 5555 and 5559, you only want to see the names of teams and

```
nba.loc[5555:5559,["team_id", "pts"]]
```

	team_id	pts
5555	FTW	83
5556	BOS	95
5557	NYK	74
5558	ROC	81
5559	SYR	86

▼ Task5: Querying the Dataset

You have seen how to access subsets of a huge dataset based on its indices, now you will select rows based on the values in your dataset's columns to query your data.

```
#create a new DataFrame that contains only games played after 2010
current_decade = nba[nba["year_id"] > 2010]
current_decade.shape
```

```
(12658, 23)
```

Question.7:

Create a new DataFrame which consists of the games played between 2000 and 2009.

```
df0 = nba.loc[(nba["year_id"] > 2000) & (nba["year_id"] < 2009)]
```

```
df0.head()
```

	gameorder	game_id	lg_id	_iscopy	year_id	date_game	seasongame	is_play
87750	43876	200010310ATL	NBA	0	2001	10/31/2000	1	
87751	43876	200010310ATL	NBA	1	2001	10/31/2000	1	
87752	43877	200010310CHI	NBA	0	2001	10/31/2000	1	
87753	43877	200010310CHI	NBA	1	2001	10/31/2000	1	
87754	43878	200010310DAL	NBA	0	2001	10/31/2000	1	

```
# save the dataframe into a csv file
df0.to_csv("df0.csv")
```

Selecting rows where a specific field is not null.

```
#selecting rows where a specific field is not null .notnull() or .notna()
games_with_notes = nba[nba["notes"].notnull()]
games_with_notes.shape
```

```
(5424, 23)
```

```
#filter your dataset and find all games where the home team's name ends with "ers".
ers = nba[nba["fran_id"].str.endswith("ers")]
ers.shape
```

```
#search for Baltimore games where both teams scored over 100 points.
#In order to see each game only once, you'll need to exclude duplicates
nba[(nba["_iscopy"] == 0) & (nba["pts"] > 100) & (nba["opp_pts"] > 100) & (nba["team_id"] ==
```

	gameorder	game_id	lg_id	_iscopy	year_id	date_game	seasongame	is_play
1726	864	194902260BLB	NBA	0	1949	2/26/1949	53	
4890	2446	195301100BLB	NBA	0	1953	1/10/1953	32	
4909	2455	195301140BLB	NBA	0	1953	1/14/1953	34	
5208	2605	195303110BLB	NBA	0	1953	3/11/1953	66	
5825	2913	195402220BLB	NBA	0	1954	2/22/1954	60	

Question.8:

Filter your dataset and find all the playoffs games where the number of points scored by both home and away teams is greater than 100, in the year 2011 and make sure you don't include duplicates (don't forget the parentheses)

```
ba[(nba["_iscopy"] == 0) & (nba["pts"] > 100) & (nba["opp_pts"] > 100) & (nba["year_id"] > 2010) & (nba["year_id"] < 2012)]
```

	gameorder	game_id	lg_id	_iscopy	year_id	date_game	seasongame	is_play
3659	56830	201010260LAL	NBA	0	2011	10/26/2010	1	
3668	56835	201010270GSW	NBA	0	2011	10/27/2010	1	
3673	56837	201010270MEM	NBA	0	2011	10/27/2010	1	

Task6: Grouping and Aggregating Your Data

You may also want to learn other features of your dataset, like the sum, mean, or average value of a group of elements. Luckily, the Pandas Python library offers grouping and aggregation functions to help you accomplish this task.

```
#Grouping - group all games for fran_id and sum their points and override the default of sort
nba.groupby("fran_id", sort=False)["pts"].sum()
```

fran_id	
Huskies	3995
Knicks	582497
Stags	20398
Falcons	3797
Capitols	22387
Celtics	626484
Steamrollers	12372
Ironmen	3674
Bombers	17793
Rebels	4474
Warriors	591224
Baltimore	37219
Jets	4482
Pistons	572758
Lakers	637444
Kings	569245
Hawks	567261
Denver	4818
Olympians	22864
Redskins	5372
Waterloo	4921
Packers	6193
Sixers	585891
Wizards	474809
Bulls	437269
Thunder	437735
Squires	91127
Stars	84940
Rockets	432504
Colonels	94435
Pacers	438288
Nuggets	445780
Spurs	453822
Spirits	85874
Sounds	75582
Floridians	49568
Nets	417809
Condors	49642
Bucks	418326
Suns	437486
Clippers	380523
Cavaliers	380416
Trailblazers	402695
Sails	30080
Jazz	363155
Mavericks	309239
Pelicans	220794
Heat	229103
Timberwolves	207693
Magic	219436
Grizzlies	157683
Raptors	158370
Hornets	84489

Name: pts, dtype: int64

```
#group by multiple columns
```

```
nba[(nba["fran_id"] == "Spurs") & (nba["year_id"] > 2010)].groupby(["year_id", "game_result"])
```

```
year_id  game_result
2011      L          25
          W          63
2012      L          20
          W          60
2013      L          30
          W          73
2014      L          27
          W          78
2015      L          31
          W          58
Name: game_id, dtype: int64
```

Question.9:

Take a look at the New York Knicks 2011-12 season (year_id: 2012). How many wins and losses did they

```
nba[(nba["fran_id"] == "Knicks") & (nba["year_id"] > 2010) & (nba["year_id"] < 2012)].groupby(["year_id", "game_result"])
```

```
year_id  game_result
2011      L          44
          W          42
Name: game_id, dtype: int64
```

The NY Knicks won 42 games and lost 44 during the 2011-2012 season.

▼ Task7: Manipulating Columns

You can add and drop columns as part of the initial data cleaning phase, or later based on the insights of your analysis.

```
#create a copy of your original DataFrame to work with
df = nba.copy()
df.shape
```

```
#define new columns based on the existing ones
df["difference"] = df.pts - df.opponent_pts
df.shape
```

```
#use an aggregation function .max() to find the largest value of your new column
```

```
dt["difference"].max()
```

```
#rename the columns of your dataset
```

```
renamed_df = df.rename(columns={"game_result": "result", "game_location": "location"})
renamed_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 126314 entries, 0 to 126313
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gameorder             126314 non-null  int64
1   game_id               126314 non-null  object
2   lg_id                126314 non-null  object
3   _iscopy              126314 non-null  int64
4   year_id              126314 non-null  int64
5   date_game            126314 non-null  object
6   seasongame           126314 non-null  int64
7   is_playoffs          126314 non-null  int64
8   team_id              126314 non-null  object
9   fran_id              126314 non-null  object
10  pts                  126314 non-null  int64
11  elo_i                126314 non-null  float64
12  elo_n                126314 non-null  float64
13  win_equiv            126314 non-null  float64
14  opp_id               126314 non-null  object
15  opp_fran             126314 non-null  object
16  opp_pts              126314 non-null  int64
17  opp_elo_i            126314 non-null  float64
18  opp_elo_n            126314 non-null  float64
19  location              126314 non-null  object
20  result               126314 non-null  object
21  forecast             126314 non-null  float64
22  notes                5424 non-null    object
23  difference            126314 non-null  int64
dtypes: float64(6), int64(8), object(10)
memory usage: 23.1+ MB
```

```
renamed_df.head(3)
```

	ts	elo_i	elo_n	win_equiv	opp_id	opp_fran	opp_pts	opp_elo_i	opp_elo_n	location
	36	1300.0	1293.28	40.29	NYK	Knicks	68	1300.00	1306.72	H
	38	1300.0	1306.72	41.71	TRH	Huskies	66	1300.00	1293.28	A
	33	1300.0	1309.65	42.01	NYK	Knicks	47	1306.72	1297.07	H

```
# save the renamed dataframe into a csv file
```

```
renamed_df.to_csv("renamed_df.csv")
```

Note that there's a new object, `renamed_df`. Like several other data manipulation methods, `.rename()` returns a new DataFrame by default.

```
#Delete unwanted columns - wont be analyzing Elo ratings here so go ahead and delete them
df.shape
```

```
elo_columns = ["elo_i", "elo_n", "opp_elo_i", "opp_elo_n"]
```

```
df.drop(elo_columns, inplace=True, axis=1)
```

```
df.shape
```

```
(126314, 20)
```

Understanding the `df.drop` function:

```
DataFrame.drop(self, labels=None, axis=0, index=None, columns=None, level=None, inplace=False, error:
```

Is translated into:

Index or column labels to drop, Whether to drop labels from the index (0 or 'index') or columns (1 or 'columns'), 'inplace=True' to make permanent changes to the dataframe.

▼ Task8: Specifying Data Types

When you create a new DataFrame, Pandas assigns a data type to each column based on its values. Sometimes is not too accurate. Choose the correct data type for your columns upfront to improve performance.

Take another look at the columns of the nba dataset:

```
#take a look at the DataFrame
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 126314 entries, 0 to 126313
Data columns (total 20 columns):
#   Column          Non-Null Count  Dtype
---  -
0   gameorder       126314 non-null  int64
1   game_id         126314 non-null  object
2   lg_id           126314 non-null  object
3   _iscopy         126314 non-null  int64
4   year_id         126314 non-null  int64
5   date_game       126314 non-null  object
6   seasongame      126314 non-null  int64
7   is_playoffs     126314 non-null  int64
8   team_id         126314 non-null  object
9   fran_id         126314 non-null  object
10  pts             126314 non-null  int64
```



```

11 win_equiv      126314 non-null float64
12 opp_id         126314 non-null object
13 opp_fran       126314 non-null object
14 opp_pts        126314 non-null int64
15 game_location  126314 non-null object
16 game_result    126314 non-null object
17 forecast       126314 non-null float64
18 notes          5424 non-null object
19 difference     126314 non-null int64
dtypes: float64(2), int64(8), object(10)
memory usage: 19.3+ MB

```

Ten of your columns have the data type object and some of these are good candidates for data type conversion.

```

# use .to_datetime() to specify all game dates as datetime objects.
df["date_game"] = pd.to_datetime(df["date_game"])

```

Similar, game_location can have only three different values. In a relational database, you would use the type enum for this column. Pandas provides the categorical data type for that same purpose.

```

#game_location column can have only three different values.
#you can see this by executing this code
df["game_location"].nunique()
df["game_location"].value_counts()

#change the data type to categorical and check it
df["game_location"] = pd.Categorical(df["game_location"])
df["game_location"].dtype

CategoricalDtype(categories=['A', 'H', 'N'], ordered=False)

```

After changing to categorical data, execute df.info. You will notice a drop in memory usage, hence improving performance.

```
df.tail(3)
```

```
game_id  game_is_playoffs  team_id  fran_id  pts  win_equiv  opp_id  opp_fran  opp_pts  game_lo
```

Question.10:

Find another column in the nba dataset that has a generic data type and convert it to a more specific

```
# convert game_result column values (W, L) into categorical type
df["game_result"].nunique()
df["game_result"].value_counts()

#change the data type to categorical and check it
df["game_result"] = pd.Categorical(df["game_result"])
df["game_result"].dtype

CategoricalDtype(categories=['L', 'W'], ordered=False)
```

▼ Task9: Cleaning the Data

▼ 9.1) Missing Values

.info() shows how many non-null values a column contains. That is very important information for you to have about your data. Null values often indicate a problem in the data-gathering process. When you inspect the dataset with nba.info() you will see that the dataset is quite neat except for the notes column which contains null values for most of its rows. This output shows that the notes column has only 5424 non-null values.

That means that over 120,000 rows of your dataset have null values in this column.

Here are a few ways to deal with null values:

```
#1st way- usually best approach is to ignore them, remove all rows with missing values
rows_without_missing_data = nba.dropna()
rows_without_missing_data.shape

(5424, 23)
```

```
#2nd way - Drop columns if they are not relevant to your analysis
# notes column is removed, rows remain unchanged
data_without_missing_columns = nba.dropna(axis=1)
data_without_missing_columns.shape

(126314, 22)
```

```
#3rd way - replace the missing values with a meaningful default value for your use case
# basically create dummy value to fill the void that would leave NAs
data_with_default_notes = nba.copy()
data_with_default_notes["notes"].fillna(value="no notes at all", inplace=True)
data_with_default_notes["notes"].describe()
```

```
count          126314
unique           232
top      no notes at all
freq          120890
Name: notes, dtype: object
```

Regarding the 1st way, that kind of data clean-up doesn't make sense for your nba dataset, because it's not a problem for a game to lack notes. But if your dataset contains a million valid records and a hundred where relevant data is missing, then dropping the incomplete records can be a reasonable solution.

▼ 9.2) Invalid Values

Use `.describe` to understand more about your dataset. This can help you identify invalid values that may throw off your analysis.

```
nba.describe()
```

	gameorder	_iscopy	year_id	seasongame	is_playoffs	pts	elo_i	
count	126314.00	126314.0	126314.00	126314.00	126314.00	126314.00	126314.00	126314.00
mean	31579.00	0.5	1988.20	43.53	0.06	102.73	1495.24	
std	18231.93	0.5	17.58	25.38	0.24	14.81	112.14	
min	1.00	0.0	1947.00	1.00	0.00	0.00	1091.64	
25%	15790.00	0.0	1975.00	22.00	0.00	93.00	1417.24	
50%	31579.00	0.5	1990.00	43.00	0.00	103.00	1500.95	
75%	47368.00	1.0	2003.00	65.00	0.00	112.00	1576.06	
max	63157.00	1.0	2015.00	108.00	1.00	186.00	1853.10	

Looking at the output you will see that the `year_id` varies between 1947 and 2015. That sounds plausible. But how can the minimum points of a game be 0.

Take a look at those games to find out if it makes sense or not.

```
#selecting the games where pts are 0
nba[nba["pts"] == 0]
```

elo_i	elo_n	win_equiv	opp_id	opp_fran	opp_pts	opp_elo_i	opp_elo_n	game_locat
1460.34	1457.45	40.41	VIR	Squires	2	1484.19	1487.08	

It seems the game was forfeited. Depending on your analysis, you may want to remove it from the dataset.

▼ 9.3) Inconsistent Values

Always check for inconsistent values. The values of the fields pts, opp_pts and game_result should be consistent with each other.

```
#check using the .empty() attribute
# here we check that teams which pts are higher than the opponent
# have indeed a winning status 'W'; empty if otherwise
# And vice-versa
nba[(nba["pts"] > nba["opp_pts"]) & (nba["game_result"] != 'W')].empty
nba[(nba["pts"] < nba["opp_pts"]) & (nba["game_result"] != 'L')].empty

True
```

Fortunately, both of these queries return an empty DataFrame. But be prepared for surprises, always check consistency.

▼ Task10: Data Visualisation

Sometimes, the numbers speak for themselves, but often a chart helps a lot with communicating your insights. Data visualizations make big and small data easier for the human brain to understand, and visualization also makes it easier to detect patterns, trends, and outliers in groups of data.

Data visualisation is one of the things that works much better in a Jupyter notebook than in a terminal. If you need help getting started, then check out **Jupyter Notebook: An Introduction**. Both


Series and DataFrame objects have a `.plot()` method, which is a wrapper around `matplotlib.pyplot.plot()`.

Visualize how many points the Knicks scored throughout the seasons.

```
#Include this line to show plots directly in the notebook
%matplotlib inline
```


```
#Visualize how many points the Knicks scored throughout the seasons
nba[nba["fran_id"] == "Knicks"].groupby("year_id")["pts"].sum().plot()
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f490a9a5090>



```
#create a bar plot to show the top 10 franchises with the most games played
nba["fran_id"].value_counts().head(10).plot(kind="bar")
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f490a939610>



Question.11 (report your answer):

11.1) Explain what the above line plot, showing how many points the Knicks scored throughout the season.

This line plot shows several peaks and valleys, with the notable sections:

- a growth trend from 1950 (most likely when data starting being recorded) to 1970, date of the first major peak (over 11000 points);
- then a descent followed by a series of lower peak until around 1995, when there is a major fall (towards 6000 points before 2000);
- Re-increase after year 2000 followed by a valley that ends with another fall around the year 2010.

11.2) Describe what the above bar plot reveals to you about the franchises with the most games played.

The bar chart shows that the Lakers are slightly ahead of the Celtics but just by a very small margin, on one hand. Furthermore, if we consider a benchmark of 5000 points, six other teams are above it. Then score drops for the last two teams (Wizards, Spurs) in the top 10.

include only the Heat's games from 2013. Then, create a plot in the same way as you've seen above).

```
# Pie plot
# where we first define a criterion to include only the Heat's games from 2013;
# Then, we counts the game results and finally plot it all with a kind 'pie'
nba[(nba["fran_id"] == "Heat") & (nba["year_id"] == 2013)]["game_result"].value_counts().plot
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f490a8aa910>
```



Double-click (or enter) to edit



▼ References:

- Pandas Docs: <https://pandas.pydata.org/pandas-docs/stable/index.html>
- Download Python: <https://www.python.org/downloads/> How to install pip: <https://www.liquidweb.com/kb/install-pip-windows/>
- Jupyter Notebook: <https://realpython.com/jupyter-notebook-introduction/>
- Install Pandas Python: https://pandas.pydata.org/pandas-docs/stable/getting_started/install.html
- Install matplotlib: <https://matplotlib.org/users/installing.html#installing-from-source>
- Visualisation with Pandas: https://pandas.pydata.org/pandasdocs/stable/user_guide/visualization.html
- Tutorial inspiration (Real Python): <https://realpython.com/pandas-python-explore-dataset/>
- Data Source: <https://fivethirtyeight.com/>
- The raw data: <https://raw.githubusercontent.com/fivethirtyeight/data/master/nba-elo/nbaallelo.csv>
- NBA Data Analysis Using Python & Machine Learning: <https://randerson112358.medium.com/nba-data-analysis-exploration-9293f311e0e8>

✓ 0s completed at 4:20 PM

